

Unit-2

- Introduction to SQL
- Intermediate and Advanced SQL

SQL History

- IBM Sequel language developed as part of System R project at the IBM San Jose Research Laboratory
- Renamed Structured Query Language (SQL)
- ANSI and ISO standard SQL:
 - SQL-86
 - SQL-89
 - SQL-92
 - SQL:1999 (language name became Y2K compliant!)
 - SQL:2003
- Commercial systems offer most, if not all, SQL-92 features, plus varying feature sets from later standards and special proprietary features.

The SQL language has several parts:

- **Data-definition language (DDL):** The SQL DDL provides commands
 - for defining relation schemas,
 - deleting relations, and
 - modifying relation schemas.
- **Data-manipulation language (DML):** The SQL DML provides the ability
 - to query information from the database and
 - to insert tuples into,
 - delete tuples from,
 - and modify tuples in the database.
- **Integrity:** The SQL DDL includes commands for specifying integrity constraints that the data stored in the database must satisfy.
 - Updates that violate integrity constraints are disallowed.

- **View definition:** The SQL DDL includes commands for defining views
- **Transaction control:** SQL includes commands for specifying the beginning and end points of transactions.
- **Embedded SQL and dynamic SQL:** Embedded and dynamic SQL define how SQL statements can be embedded within general-purpose programming languages, such as C, C++, and Java
- **Authorization:** The SQL DDL includes commands for specifying access rights to relations and views.

SQL Data Definition

Allows the specification of not only a set of relations but also information about each relation, including:

- The schema for each relation.
- The domain of values associated with each attribute.
- Integrity constraints
- The set of indices to be maintained for each relations.
- Security and authorization information for each relation.
- The physical storage structure of each relation on disk.

Basic Types

- **char(n)**: Fixed length character string, with user-specified length n .
- **varchar(n)**: Variable length character strings, with user-specified maximum length n .
- **int**: Integer (a finite subset of the integers that is machine-dependent).
- **smallint**: Small integer (a machine-dependent subset of the integer domain type).
- **number(p,d)**: Fixed point number, with user-specified precision of p digits, with n digits to the right of decimal point.
- **float(n)**: Floating point number, with user-specified precision of at least n digits.
- **date**: A calendar date containing a (four-digit) year, month, and day of the month.

Create Table

- **Creating a Database**
 - create database database-name;
 - create database stud_db;
- An SQL relation is defined using the **create table** command:
- **create table** $r(A_1 D_1, A_2 D_2, \dots, A_n D_n, (\text{integrity-constraint}_1)^{n_1}, \dots, (\text{integrity-constraint}_k)^{n_k});$
 - r is the name of the relation
 - each A_i is an attribute name in the schema of relation r
 - D_i is the data type of values in the domain of attribute A_i
- Example: **Schema:branch(branch_name:string,branch_city:string, assets:int);**
 - **create table** *branch*
(*branch_name* char(15) **not null**,
branch_city char(30),
assets integer);

Integrity Constraints in Create Table

- **not null**
- **primary key** (A_1, \dots, A_n)
 - Example: Declare *branch_name* as the primary key for *branch*
 - **create table** *branch*
(*branch_name* char(15),
branch_city char(30),
assets integer,
primary key (*branch_name*))
- primary key ($A_{j1}, A_{j2}, \dots, A_{jm}$): The primary-key specification says that attributes $A_{j1}, A_{j2}, \dots, A_{jm}$ form the primary key for the relation. The primary-key attributes
- **primary key** declaration on an attribute automatically ensures **not null** in SQL-92 onwards

foreign key ($A_{k1}, A_{k2}, \dots, A_{kn}$) **references** s:

The foreign key specification says that the values of attributes ($A_{k1}, A_{k2}, \dots, A_{kn}$) for any tuple in the relation must correspond to values of the primary key attributes of some tuple in relation s.

- **Drop command**
- To remove a relation from an SQL database, we use the drop table command.
- The drop table command deletes all information about the dropped relation from the database.
 - drop table r;
 - Ex:
 - Drop table student;
- **Truncate command**
- It is used to delete all the tuples from the table.
 - TRUNCATE TABLE table_name;
 - Ex:
 - Truncate table student;

Alter command

- The **alter table** command is used to add attributes to an existing relation:
- **alter table r add A D**
 - where A is the name of the attribute to be added to relation r and D is the domain of A.
 - All tuples in the relation are assigned *null* as the value for the new attribute.
 - Example: alter table Student add(address char(20));
 -
- To change the **data type of a column** in a table, use the following syntax:
 - ALTER TABLE r MODIFY COLUMN A D;
 - Example: alter table Student modify(address varchar(30));
- To **Rename a column**
- alter table table-name rename old-column-name to column-name;
- alter table Student rename address to Location;
-

- The **alter table** command can also be used to **drop attributes** of a relation:

alter table r drop A

where A is the name of an attribute of relation r

Example: alter table Student drop(address);

Alter table student drop constraint myprimarykey;

- To change the **name of a table** or relation
 - ALTER TABLE table_name RENAME TO new_table_name;
 - Ex: alter table student rename to stud

Adding constraints

alter table table-name add constraint myprimarykey primarykey (col1,col2..);

Schema

student(sid: number, name: string, branch: string, section: string, age: number, cgpa: number(2,1))

Data Manipulation Language(Modification of the Database)

- **insertion**
- To insert data/row into a relation,
 - we either specify a tuple to be inserted
 - 1) Insert into r values(valu1,value2,.....valuen);
 - Ex: insert into student values('561','ABC','CSE','B',20,9.7);
 - 2) INSERT into table_name(column1,column2,...) values(data1,data2,.....);
 - Ex: INSERT into Student(sid,name,branch,age,section,cgpa) values('562','Ravi','CSE',20,'B',9.7);
 - or
 - write a query whose result is a set of tuples to be inserted
 - Ex: insert into r select attr1,attr2,attr3 from r2 where condition;

- **UPDATE** command is used to update the value of a specified attribute of a row or rows

- UPDATE table_name
set column_name = value
where condition;
 - EX:
UPDATE Student
set name='Abhi',age=17
where sid='565';

- **Delete** command is used to delete the rows from a table

- DELETE from table_name
where condition;
- Ex:
 - DELETE from Student
 - where sid=103;

Basic Query Structure

SELECT [DISTINCT] *target-list*
FROM *relation-list*
WHERE *qualification*

- **target-list** :A list of attributes of relations
- **relation-list** A list of relation names
- **qualification**:in *relation-list* *qualification* Comparisons combined using AND, OR, and NOT
 - Attr op const or Attr1 op Attr2,
 - where op is one of relational operator
- **DISTINCT** is an optional keyword indicating that the answer should not contain duplicates

- Semantics of an SQL query defined in terms of the following conceptual evaluation strategy:

- Compute the cross-product of *relation-list*.
- Discard resulting tuples if they fail *qualifications*.
- Delete attributes that are not in *target-list*.
- If DISTINCT is specified, eliminate duplicate rows.

Examples:

- Display all the student details
Select *
from student ;
- Display distinct names of students
Select distinct name
from student
- Display B Section students
 - Select *
 - from student
 - where section='B';

- Display htno,name and cgpa of B Section student
 - Select htno,name,age,cgpa
 - from student
 - where section='B';

Ordering the Display of Tuples

*SELECT [DISTINCT] target-list
FROM relation-list
WHERE qualification
Order by column_name sort_order;*

- *sort_order* may be either **desc** or **asc**

- Display the students order by their names
 - Select * from student orderby name;
- Display id and names of students order by their cgpa
- Select sid,name from student orderby cgpa;
- Display B section students order by their cgpa

```
SELECT * FROM student WHERE section='B' order by cgpa
```

String Operations

- SQL includes a string-matching operator for comparisons on character strings.
- The operator “like” uses patterns that are described using two special characters:
- percent (%). The % character matches any substring(‘%’ stands for 0 or more characters)
- underscore (_). The _ character matches any character.(‘_’ stands for any one character)

- Find the ages of students whose name begins and ends with B and has at least three characters.
- SELECT age
FROM Student
WHERE name LIKE `B_ %B`

Aggregate Functions

SQL supports five aggregate operations, which can be applied on any column, say A, of a relation:

- COUNT ([DISTINCT] A): The number of (unique) values in the A column.
 - SELECT COUNT (*) FROM students
- SUM ([DISTINCT] A): The sum of all (unique) values in the A column.
 - Select sum(cost) from Products;

- AVG ([DISTINCT] A): The average of all (unique) values in the A column.
 - Ex: Find the average age of all students
 - SELECT AVG (age) FROM student
- MAX (A): The maximum value in the A column.
 - Find the maximum cgpa of student
 - SELECT MAX(cgpa) FROM student
- MIN (A): The minimum value in the A column.
 - Find the minimum cgpa of student
 - SELECT MIN(cgpa) FROM student

Sailors(sid: integer, sname: string, rating: integer, age: real)

Boats(bid: integer, bname: string, color: string)

Reserves(sid: integer, bid: integer, day: date)

```
CREATE TABLE sailors ( sid integer, sname varchar(32), rating integer, age real, PRIMARY KEY (sid) );
```

```
CREATE TABLE boats( bid integer, bname string, color string, PRIMARY KEY (bid) );
```

```
CREATE TABLE reserves ( sid integer , bid integer , day date,
PRIMARY KEY (sid, bid, day),
FOREIGN KEY (sid) REFERENCES sailors(sid),
FOREIGN KEY (bid) REFERENCES boats(bid) );
```

Sailors

<i>sid</i>	<i>sname</i>	<i>rating</i>	<i>age</i>
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

Reserves

<i>sid</i>	<i>bid</i>	<i>day</i>
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

Boats

<i>bid</i>	<i>bname</i>	<i>color</i>
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

- Find the names and ages of all sailors.

SELECT DISTINCT S.sname, S.age FROM Sailors S

With DISTINCT

<i>sname</i>	<i>age</i>
Dustin	45.0
Brutus	33.0
Lubber	55.5
Andy	25.5
Rusty	35.0
Horatio	35.0
Zorba	16.0
Art	25.5
Bob	63.5

Without DISTINCT

<i>sname</i>	<i>age</i>
Dustin	45.0
Brutus	33.0
Lubber	55.5
Andy	25.5
Rusty	35.0
Horatio	35.0
Zorba	16.0
Horatio	35.0
Art	25.5
Bob	63.5

- Find all sailors with a rating above 7.

SELECT S.sid, S.sname, S.rating, S.age

FROM Sailors AS S

WHERE S.rating > 7

- Find the names of sailors who have reserved boat number 103

SELECT S.sname

FROM Sailors S, Reserves R

WHERE S.sid = R.sid AND R.bid=103

OR

SELECT sname

FROM Sailors, Reserves

WHERE Sailors.sid=Reserves.sid AND bid=103

Sailors S

<i>sid</i>	<i>sname</i>	<i>rating</i>	<i>age</i>
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

Reserves R

<i>sid</i>	<i>bid</i>	<i>day</i>
22	101	10/10/96
58	103	11/12/96

S x R

<i>sid</i>	<i>sname</i>	<i>rating</i>	<i>age</i>	<i>sid</i>	<i>bid</i>	<i>day</i>
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

Result

sname

rusty

- Find the sids of sailors who have reserved a red boat

```
Select R.sid
FROM Boats B, Reserves R
WHERE B.bid = R.bid AND B.color = 'red'
```

- Find the names of sailors who have reserved a red boat.

```
SELECT S.sname
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'red'
```

- Find the colors of boats reserved by Lubber

```
SELECT B.color
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND S.sname = 'Lubber'
```

- Find the names of sailors who have reserved at least one boat.

```
SELECT S.sname
FROM Sailors S, Reserves R
WHERE S.sid = R.sid
```

Expressions and Strings in the SELECT Command

- Each item in a **select-list** can be of the form

- *expression AS column name*,
- where *expression* is any arithmetic or string expression over column names and constants

- Compute the increments ratings of persons who have sailed two different boats on the same day.

```
SELECT S.sname, S.rating+1 AS rating
FROM Sailors S, Reserves R1, Reserves R2
Where S.sid=R1.sid AND S.sid=R2.sid AND R1.day=R2.day AND R1.bid<>R2.bid;
```

- Find the ages of sailors whose name begins and ends with B and has at least three characters.

- SELECT S.age
- FROM Sailors S
- WHERE S.sname LIKE 'B_ %B'

- Find the names of sailors who have reserved a red or a green boat.

```
SELECT S.sname
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
AND (B.color = 'red' OR B.color = 'green')
```

- Find the names of sailors who have reserved both a red and a green boat.

```
SELECT S.sname
FROM Sailors S, Reserves R1, Boats B1, Reserves R2, Boats B2
WHERE S.sid = R1.sid AND R1.bid = B1.bid
AND S.sid = R2.sid AND R2.bid = B2.bid
AND B1.color='red' AND B2.color = 'green'
```

Set Operations(UNION, INTERSECT, AND EXCEPT)

- SQL provides three set-manipulation constructs that extend the basic query form presented earlier

UNION
INTERSECT
EXCEPT

- Find the names of sailors who have reserved a red or a green boat

```
SELECT S.sname
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'red'
UNION
SELECT S2.sname
FROM Sailors S2, Boats B2, Reserves R2
WHERE S2.sid = R2.sid AND R2.bid = B2.bid AND B2.color = 'green'
```

- Find the names of sailors who have reserved both a red and a green boat

```
SELECT S.sname
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'red'
INTERSECT
SELECT S2.sname
FROM Sailors S2, Boats B2, Reserves R2
WHERE S2.sid = R2.sid AND R2.bid = B2.bid AND B2.color = 'green'
```

- Find the sids of all sailors who have reserved red boats but not green boats.

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid = B.bid AND B.color = 'red'
EXCEPT
SELECT R2.sid
FROM Boats B2, Reserves R2
WHERE R2.bid = B2.bid AND B2.color = 'green'
```

- Find all sids of sailors who have a rating of 10 or have reserved boat 104

```
SELECT S.sid
FROM Sailors S
WHERE S.rating = 10
UNION
SELECT R.sid
FROM Reserves R
WHERE R.bid = 104
```

- SQL also provides other set operations:
 - IN -to check if an element is in a given set
 - op ANY, op ALL -to compare a value with the elements in a given set, using comparison operator op
 - EXISTS - to check if a set is empty

NESTED QUERIES

- A **nested query** is a **query** that has another query embedded within it
- the embedded query is called a **sub query**.
- A sub query typically appears within the WHERE clause of a query
- Sub queries can sometimes appear in the FROM clause or the HAVING clause

- Find the names of sailors who have reserved boat 103.

```
SELECT S.sname
FROM Sailors S
WHERE S.sid IN ( SELECT R.sid
                  FROM Reserves R
                  WHERE R.bid = 103 )
```

- Find the names of sailors who have reserved a red boat.

```
SELECT S.sname
FROM Sailors S
WHERE S.sid IN ( SELECT R.sid
                  FROM Reserves R
                  WHERE R.bid IN ( SELECT B.bid
                                  FROM Boats B
                                  WHERE B.color = 'red' ))
```

- *Find the names of sailors who have not reserved a red boat.*

[illegible]

Correlated Nested Queries

- In the nested queries that we have seen thus far, the inner sub query has been completely independent of the outer query
- In general the inner sub query could depend on the row that is currently being examined in the outer query

- Find the names of sailors who have reserved boat number 103

```
SELECT S.sname  
FROM Sailors S  
WHERE EXISTS ( SELECT *  
                FROM Reserves R  
                WHERE R.bid = 103  
                AND R.sid = S.sid )
```

- The EXISTS operator is another set comparison operator, such as IN
 - It allows us to test whether a set is nonempty
 - for each Sailor row *S*, we test whether the set of Reserves rows *R* such that *R.bid = 103 AND S.sid = R.sid* is nonempty.
 - If so, sailor *S* has reserved boat 103, and we retrieve the name

- The sub query clearly depends on the current row *S* and must be re-evaluated for each row in *Sailors*.
- The occurrence of *S* in the sub query (in the form of the literal *S.sid*) is called a *correlation*, and such queries are called *correlated queries*
- using NOT EXISTS instead of EXISTS, we can compute the names of sailors who have not reserved a red boat

Set-Comparison Operators

- SQL also supports op ANY and op ALL,
 - where op is one of the arithmetic comparison operators
 - <; <=;; =; >; >=;; >
 - *SOME* is also available, but it is just a synonym for ANY

- Find sailors whose rating is better than some sailor called Horatio

```
SELECT S.sid
FROM Sailors S
WHERE S.rating > ANY ( SELECT S2.rating
                      FROM Sailors S2
                      WHERE S2.sname = 'Horatio' );
```

- Find sailors whose rating is better than every sailor called Horatio.

```
SELECT S.sid
FROM Sailors S
WHERE S.rating > ALL ( SELECT S2.rating
                     FROM Sailors S2
                     WHERE S2.sname = 'Horatio' );
```

- Find the sailors with the highest rating.

```
SELECT S.sid
FROM Sailors S
WHERE S.rating >= ALL ( SELECT S2.rating
                      FROM Sailors S2 )
```

- Find the names of sailors who have reserved both a red boat and a green boat

Select s.sname

from sailors s

Where s.sid IN((select R.sid

from Boats B, Reserves R

where R.bid=B.bid AND B.color='red')

INTERSECT

(select R.sid

from Boats B2, Reserves R2

where R2.bid=B2.bid AND B2.color='green'));

- Find the Names of sailors who have reserved all boats

Select s.sname

from sailors s

Where NOT EXISTS(select B.bid

from Boats B

where NOT EXISTS(select R.bid

from Reserves R

where R.bid=B.bid

AND R.sid=S.sid));

- Count the number of sailors.

- SELECT COUNT (*)

- FROM Sailors S;

- Count the number of different sailor names.

- SELECT COUNT (DISTINCT S.sname)

- FROM Sailors S

- *Find the name and age of the oldest sailor.*

- `SELECT S.sname, MAX (S.age)`
- `FROM Sailors S`

- if the SELECT clause uses an aggregate operation, then it must use *only aggregate operations unless* the query contains a GROUP BY clause

- `SELECT S.sname, S.age`
- `FROM Sailors S`
- `WHERE (SELECT MAX (S2.age)`
- `FROM Sailors S2) = S.age`

- *Find the names of sailors who are older than the oldest sailor with a rating of 10.*

- `SELECT S.sname`
- `FROM Sailors S`
- `WHERE S.age > (SELECT MAX (S2.age)`
`FROM Sailors S2`
`WHERE S2.rating = 10)`

OR

- `SELECT S.sname`
- `FROM Sailors S`
- `WHERE S.age > ALL (SELECT S2.age`
`FROM Sailors S2`
`WHERE S2.rating = 10)`

The GROUP BY and HAVING Clauses

- Often we want to apply aggregate operations to each of a number of groups of rows in a relation,
 - where the number of groups depends on the relation instance
- `SELECT [DISTINCT] select-list`
- `FROM from-list`
- `WHERE qualification`
- `GROUP BY grouping-list`
- `HAVING group-qualification`

- The select-list in the SELECT clause consists of
 - (1) a list of column names and
 - (2) a list of terms having the form *aggop (column-name) AS new-name*.
- Every column that appears in (1) must also appear in grouping-list.
 - The reason is that each row in the result of the query corresponds to one *group*, *which is a* collection of rows that agree on the values of columns in grouping-list.
 - If a column appears in list (1), but not in grouping-list, it is not clear what value should be assigned to it in an answer row.

- The expressions appearing in the group-qualification in the HAVING clause must have a *single value per group*.
 - a column appearing in the group-qualification must appear as the argument to an aggregation operator, or it must also appear in grouping-list.
- If the GROUP BY clause is omitted, the entire table is regarded as a single group.

- Find the age of the youngest sailor for each rating level.
- SELECT S.rating, MIN (S.age)
- FROM Sailors S
- GROUP BY S.rating

- *Find the age of the youngest sailor who is eligible to vote for each rating level with at least two such sailors.*
 - SELECT S.rating, MIN (S.age) AS minage
 - FROM Sailors S
 - WHERE S.age >= 18
 - GROUP BY S.rating
 - HAVING COUNT (*) > 1

- *For each red boat, find the number of reservations for this boat.*

- SELECT B.bid, COUNT (*) AS sailorcount
- FROM Boats B, Reserves R
- WHERE R.bid = B.bid
- GROUP BY B.bid
- HAVING B.color = 'red'

- *Find the average age of sailors for each rating level that has at least two sailors.*

- SELECT S.rating, AVG (S.age) AS avgage
- FROM Sailors S
- GROUP BY S.rating
- HAVING COUNT (*) > 1

- *Find those ratings for which the average age of sailors is the minimum over all ratings.*

```
SELECT Temp.rating, Temp.avgage
FROM ( SELECT S.rating, AVG (S.age) AS avgage,
FROM Sailors S
GROUP BY S.rating) AS Temp
WHERE Temp.avgage = ( SELECT MIN (Temp.avgage) FROM Temp )
```

NULL VALUES

- SQL provides a special column value called null .
- We use null when the column value is either unknown or inapplicable.
- The presence of null values complicates many issues, and we consider the impact of null values

- SQL also provides a special comparison operator IS NULL to test whether a column value is *null*;

- IS NOT NULL

- Example: select loan_number from loan where amount is null