

UNIT-3

- Formal Relational Query Languages
- Relational Database Design

Formal Relational Query Languages

- Relational Algebra
- Relational Calculus
 - Domain Relational Calculus
 - Tuple Relational Calculus

Relational Query Languages

- A major strength of the relational model: supports simple, powerful *querying* of data.
- Queries can be written intuitively, and the DBMS is responsible for efficient evaluation.
 - The key: precise semantics for relational queries.
 - Allows the optimizer to extensively re-order operations, and still ensure that the answer does not change.
- Query languages can be categorized into 2 types
 - 1.Procedural Language
 - 2.Non-procedural Language

Relational Query Languages(contd.)

- **Procedural Language:** user instructs the system to perform a sequence of operations on the database to compute the desired result.
 - Relational Algebra is one of the Procedural Language
- **Non Procedural Language:** User describes the desired information without giving a specific procedure for obtaining that information
 - Relational calculus is one of the Non Procedural Language
 - RC again categorized into TRC,DRC

Relational Algebra

- Relational Algebra is formal description of how relational database operates.
- It is a procedural query language, i.e. user must define both “how” and “what” to retrieve.
- It consists of a set of operators that consume either one or two relations as input.
 - An operator produces one relation as its output.

Relational Algebra

- The fundamental operations are :
 - Select
 - Project
 - Union
 - Set difference
 - Cartesian Product
 - Rename
 - Set intersection
 - Division
 - Assignment
 - Natural join

Algebra Operations

- Unary Operations - operate on one relation.
 - select, project and rename operators.
- Binary Operations - operate on pairs of relations.
 - union, set difference, division, Cartesian product, equality join, natural join, join and semi-join operators.

Select Operator

- The **Select** operator selects tuples that satisfies a predicate; e.g. retrieve the employees whose salary is 30,000
 - $\sigma_{Salary = 30,000}(Employee)$
- Conditions in Selection:
 - Simple Condition: $(attribute)(comparison)(attribute)$
 - $(attribute)(comparison)(constant)$
 - Comparison: $=, \neq, \leq, \geq, <, >$
 - Condition: *combination of simple conditions with AND, OR, NOT*
-

Select Operator Example

Person

| Name | Age | Weight |
|--------|-----|--------|
| Harry | 34 | 80 |
| Sally | 28 | 64 |
| George | 29 | 70 |
| Helena | 54 | 54 |
| Peter | 34 | 80 |

$\sigma_{\text{Age} \geq 34}(\text{Person})$

| Name | Age | Weight |
|--------|-----|--------|
| Harry | 34 | 80 |
| Helena | 54 | 54 |
| Peter | 34 | 80 |

$\sigma_{\text{Age}=\text{Weight}}(\text{Person})$

| Name | Age | Weight |
|--------|-----|--------|
| Helena | 54 | 54 |

Project Operator

- **Project (Π)** retrieves a column. Duplication is not permitted.
- e.g., name of employees:
 - $\Pi_{\text{name}}(\text{Employee})$

Employee

| Name | Age | Salary |
|--------|-----|--------|
| Harry | 34 | 80,000 |
| Sally | 28 | 90,000 |
| George | 29 | 70,000 |
| Helena | 54 | 54,280 |
| Peter | 34 | 40,000 |

$\Pi_{\text{name}}(\text{Employee})$

| Name |
|--------|
| Harry |
| Sally |
| George |
| Helena |
| Peter |

Composite Example

Eg: Name of employees earning more than 80,000:

$\Pi_{\text{name}}(\sigma_{\text{Salary} > 80,000}(\text{Employee}))$

Employee

| Name | Age | Salary |
|--------|-----|--------|
| Harry | 34 | 80,000 |
| Sally | 28 | 90,000 |
| George | 29 | 70,000 |
| Helena | 54 | 54,280 |
| Peter | 34 | 40,000 |

$\sigma_{\text{Salary} > 80,000}(\text{Employee})$

| Name | Age | Salary |
|-------|-----|--------|
| Sally | 28 | 90,000 |

$\Pi_{\text{name}}(\sigma_{\text{Salary} > 80,000}(\text{Employee}))$

| Name |
|-------|
| Sally |

Cartesian Product

- In mathematics, it is a set of all pairs of elements (x, y) that can be constructed from given sets, X and Y, such that x belongs to X and y to Y.
- It defines a relation that is the concatenation of every tuple of relation R with every tuple of relation S.

Cartesian Product Example

| Person | | | City | |
|--------|-----|--------|----------|--|
| Name | Age | Weight | City | |
| Harry | 34 | 80 | San Jose | |
| Sally | 28 | 64 | San Jose | |
| George | 29 | 70 | San Jose | |
| | | | Austin | |

| Person X City | | | |
|---------------|-----|--------|----------|
| Name | Age | Weight | City |
| Harry | 34 | 80 | San Jose |
| Sally | 28 | 64 | San Jose |
| George | 29 | 70 | San Jose |
| Harry | 34 | 80 | Austin |
| Sally | 28 | 64 | Austin |
| George | 29 | 70 | Austin |

Example

Cartesian Product (\times)

- Allows us to combine information from any two relations. Cartesian product of relations r_1 and r_2 are represented as $r_1 \times r_2$.

| r_1 | | r_2 | | | r | | | | |
|-------|-------|-------|-------|-------|-------|---------|---------|-------|-------|
| A | B | B | C | D | A | $r_1.B$ | $r_2.B$ | C | D |
| a_1 | b_1 | b_1 | c_1 | d_1 | a_1 | b_1 | b_1 | c_1 | d_1 |
| a_2 | b_2 | b_2 | c_2 | d_2 | a_1 | b_1 | b_2 | c_2 | d_2 |
| | | b_3 | c_3 | d_3 | a_2 | b_2 | b_3 | c_3 | d_3 |
| | | | | | a_2 | b_2 | b_2 | c_2 | d_2 |
| | | | | | a_2 | b_2 | b_3 | c_3 | d_3 |

Introduction to databases

IL-11

AKN/DB

Rename Operator

- In relational algebra, a **rename** is a unary operation written as $\rho_{a/b}(R)$
 - where a and b are attribute names and R is a relation
- The result is identical to R except that the b field in all tuples is renamed to an a field.
- Example
 - $\rho_{\text{employeename/Name}}(\text{Emp})$
 - Changes the name of column 'Name' to 'EmployeeName' in Emp table

Rename Operator Example

| Employee | | $\rho_{\text{EmployeeName/Name}}(\text{Employee})$ | |
|----------|--------|--|--------|
| Name | Salary | EmployeeName | Salary |
| Harry | 80,000 | Harry | 80,000 |
| Sally | 90,000 | Sally | 90,000 |
| George | 70,000 | George | 70,000 |
| Helena | 54,280 | Helena | 54,280 |
| Peter | 40,000 | Peter | 40,000 |

Union Operator

- The **union** operation is denoted **U** as in set theory.
- It returns the union (set union) of two compatible relations.
- For a union operation $r \cup s$ to be legal, we require that,
 - r and s must have the same number of attributes.
 - The domains of the corresponding attributes must be the same.
- As in all set operations, duplicates are eliminated.

Union Operator Example

| FN | LN |
|---------|-------|
| Susan | Yao |
| Ramesh | Shah |
| Barbara | Jones |
| Amy | Ford |
| Jimmy | Wang |

| FN | LN |
|---------|---------|
| John | Smith |
| Ricardo | Brown |
| Susan | Yao |
| Francis | Johnson |
| Ramesh | Shah |

Student U Professor

| FN | LN |
|---------|---------|
| Susan | Yao |
| Ramesh | Shah |
| Barbara | Jones |
| Amy | Ford |
| Jimmy | Wang |
| John | Smith |
| Ricardo | Brown |
| Francis | Johnson |

Intersection Operator

- Denoted as \cap .
- For relations R and S , intersection is $R \cap S$.
 - Defines a relation consisting of the set of all tuples that are in both R and S .
 - R and S must be union-compatible.
- Expressed using basic operations:
 - $R \cap S = R - (R - S)$

Intersection Operator Example

| FN | LN |
|---------|-------|
| Susan | Yao |
| Ramesh | Shah |
| Barbara | Jones |
| Amy | Ford |
| Jimmy | Wang |

| FN | LN |
|---------|---------|
| John | Smith |
| Ricardo | Brown |
| Susan | Yao |
| Francis | Johnson |
| Ramesh | Shah |

Student \cap Professor

| FN | LN |
|--------|------|
| Susan | Yao |
| Ramesh | Shah |

Set Difference Operator

- For relations R and S ,
 - Set difference $R - S$, defines a relation consisting of the tuples that are in relation R , but not in S .
 - Set difference $S - R$, defines a relation consisting of the tuples that are in relation S , but not in R .

Set Difference Operator Example

Student

| FN | LN |
|---------|-------|
| Susan | Yao |
| Ramesh | Shah |
| Barbara | Jones |
| Amy | Ford |
| Jimmy | Wang |

Professor

| FN | LN |
|---------|---------|
| John | Smith |
| Ricardo | Brown |
| Susan | Yao |
| Francis | Johnson |
| Ramesh | Shah |

Professor - Student

| FN | LN |
|---------|---------|
| John | Smith |
| Ricardo | Brown |
| Francis | Johnson |

Student - Professor

| FN | LN |
|---------|-------|
| Barbara | Jones |
| Amy | Ford |
| Jimmy | Wang |

Division Operator

- The division operator takes as input two relations, called the dividend relation (r on schema R) and the divisor relation (s on schema S) such that all the attributes in S also appear in R and S is not empty.
- The output of the division operation is a relation on schema R with all the attributes common with S .

Division Operator Example

Completed

| Student | Task |
|---------|-----------|
| Fred | Database1 |
| Fred | Database2 |
| Fred | Compiler1 |
| Eugene | Database1 |
| Sara | Database1 |
| Sara | Database2 |
| Eugene | Compiler1 |

DBProject

| Task |
|-----------|
| Database1 |
| Database2 |

Completed / DBProject

| Student |
|---------|
| Fred |
| Sara |

DIVISION

MAIN

| A | B |
|----|----|
| A1 | B1 |
| A3 | B2 |
| A2 | B3 |
| A2 | B1 |
| A1 | B3 |
| A3 | B3 |
| A3 | B1 |

SUB1

| B |
|----|
| B1 |
| B3 |

SUB2

| B |
|----|
| B2 |

SUB3

| B |
|----|
| B2 |
| B3 |

ASSIGNMENT

$R1 \leftarrow \text{MAIN} \div \text{SUB1}$

$R2 \leftarrow \text{MAIN} \div \text{SUB2}$

$R3 \leftarrow \text{MAIN} \div \text{SUB3}$

RESULTING TABLES

R1

| A |
|----|
| A1 |
| A2 |
| A3 |

R2

| A |
|----|
| A3 |

R3

| A |
|----|
| A3 |

12

Joins

- It is used to combine information from two or more relations
- It is defined as a cross product followed by selections and projections
- Joins are of two types
 - Inner Joins
 - Outer Joins
- Inner Joins:** An inner join includes only those tuples with matching attributes and the rest are discarded in the resulting relation.
 - Condition Join, Equijoin, and Natural Join are called inner joins.

Joins(contd..)

- Outer Joins:** To include all the tuples from the participating relations in the resulting relation. (it allows NULL values)
- There are three kinds of outer joins –
 - left outer join,
 - right outer join, and
 - full outer join.

Notation

| Operation | My HTML | Symbol |
|--------------|--------------|--------------|
| Projection | PROJECT | π |
| Selection | SELECT | σ |
| Renaming | RENAME | ρ |
| Union | UNION | \cup |
| Intersection | INTERSECTION | \cap |
| Assignment | \leftarrow | \leftarrow |

| Operation | My HTML | Symbol |
|-------------------|------------------|------------------|
| Cartesian product | X | \times |
| Join | JOIN | \bowtie |
| Left outer join | LEFT OUTER JOIN | \ltimes |
| Right outer join | RIGHT OUTER JOIN | \rtimes |
| Full outer join | FULL OUTER JOIN | $\ltimes\rtimes$ |
| Semijoin | SEMIJOIN | \ltimes |

Condition Join Operator

- The join operation accepts a join condition c and a pair of relation instances as arguments and returns a relation instance
- $R \bowtie_c S = \sigma_c(R \times S)$

Example: Bag Theta-Join

R(

| A, | B |
|----|---|
| 1 | 2 |
| 5 | 6 |
| 1 | 2 |

)

S(

| B, | C |
|----|---|
| 3 | 4 |
| 7 | 8 |

)

| $R \text{ JOIN}_{R.B < S.B} S =$ | <table><tr><th>A</th><th>R.B</th><th>S.B</th><th>C</th></tr><tr><td>1</td><td>2</td><td>3</td><td>4</td></tr><tr><td>1</td><td>2</td><td>7</td><td>8</td></tr><tr><td>5</td><td>6</td><td>7</td><td>8</td></tr><tr><td>1</td><td>2</td><td>3</td><td>4</td></tr><tr><td>1</td><td>2</td><td>7</td><td>8</td></tr></table> | A | R.B | S.B | C | 1 | 2 | 3 | 4 | 1 | 2 | 7 | 8 | 5 | 6 | 7 | 8 | 1 | 2 | 3 | 4 | 1 | 2 | 7 | 8 |
|----------------------------------|---|-----|-----|-----|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| A | R.B | S.B | C | | | | | | | | | | | | | | | | | | | | | | |
| 1 | 2 | 3 | 4 | | | | | | | | | | | | | | | | | | | | | | |
| 1 | 2 | 7 | 8 | | | | | | | | | | | | | | | | | | | | | | |
| 5 | 6 | 7 | 8 | | | | | | | | | | | | | | | | | | | | | | |
| 1 | 2 | 3 | 4 | | | | | | | | | | | | | | | | | | | | | | |
| 1 | 2 | 7 | 8 | | | | | | | | | | | | | | | | | | | | | | |

57

Natural Join Operator

- Natural join is a dyadic operator that is written as $R \bowtie S$ where R and S are relations.
- The result of the natural join is the set of all combinations of tuples in R and S that are equal on their common attribute names.

Natural Join Example

For an example, consider the tables *Employee* and *Dept* and their natural join:

Employee

| Name | EmpID | DeptName |
|---------|-------|----------|
| Harry | 3415 | Finance |
| Sally | 2241 | Sales |
| George | 3401 | Finance |
| Harriet | 2202 | Sales |

Dept

| DeptName | Mgr |
|------------|---------|
| Finance | George |
| Sales | Harriet |
| Production | Charles |

Employee \bowtie Dept

| Name | EmpID | DeptName | Mgr |
|---------|-------|----------|---------|
| Harry | 3415 | Finance | George |
| Sally | 2241 | Sales | Harriet |
| George | 3401 | Finance | George |
| Harriet | 2202 | Sales | Harriet |

Semijoin Operator

- The semijoin is joining similar to the natural join and written as $R \bowtie S$ where R and S are relations.
- The result of the semijoin is only the set of all tuples in R for which there is a tuple in S that is equal on their common attribute names.

Semijoin Example

For an example consider the tables *Employee* and *Dept* and their semi join:

| Name | EmpID | DeptName |
|---------|-------|----------|
| Harry | 3415 | Finance |
| Sally | 2241 | Sales |
| George | 3401 | Finance |
| Harriet | 2202 | Sales |

| DeptName | Mgr |
|------------|---------|
| Sales | Harriet |
| Production | Charles |

Employee \bowtie Dept

| Name | EmpID | DeptName |
|---------|-------|----------|
| Sally | 2241 | Sales |
| Harriet | 2202 | Sales |

Outer joins

- **Left outer join**
- The left outer join is written as $R =X S$ where R and S are relations.
- The result of the left outer join is the set of all combinations of tuples in R and S that are equal on their common attribute names, in addition to tuples in R that have no matching tuples in S .

Left Outerjoin Example

For an example consider the tables *Employee* and *Dept* and their left outer join:

| Name | EmpID | DeptName |
|---------|-------|----------|
| Harry | 3415 | Finance |
| Sally | 2241 | Sales |
| George | 3401 | Finance |
| Harriet | 2202 | Sales |

| DeptName | Mgr |
|----------|---------|
| Sales | Harriet |

Employee $=X$ Dept

| Name | EmpID | DeptName | Mgr |
|---------|-------|----------|---------|
| Harry | 3415 | Finance | NULL |
| Sally | 2241 | Sales | Harriet |
| George | 3401 | Finance | NULL |
| Harriet | 2202 | Sales | Harriet |

Outer joins(contd..)

- **Right outer join**
- The right outer join is written as $R \bowtie_r S$ where R and S are relations.
- The result of the right outer join is the set of all combinations of tuples in R and S that are equal on their common attribute names, in addition to tuples in S that have no matching tuples in R .

Right Outerjoin Example

For an example consider the tables *Employee* and *Dept* and their right outer join:

| Employee | | | Dept | |
|----------|-------|----------|------------|---------|
| Name | EmpID | DeptName | DeptName | Mgr |
| Harry | 3415 | Finance | Sales | Harriet |
| Sally | 2241 | Sales | Production | Charles |
| George | 3401 | Finance | | |
| Harriet | 2202 | Sales | | |

Employee \bowtie_r Dept

| Name | EmpID | DeptName | Mgr |
|---------|-------|------------|---------|
| Sally | 2241 | Sales | Harriet |
| Harriet | 2202 | Sales | Harriet |
| NULL | NULL | Production | Charles |

Full Outer join Example

The **outer join** or **full outer join** in effect combines the results of the left and right outer joins.

| Employee | | | Employee \bowtie_{full} Dept | | | |
|----------|-------|----------|--------------------------------|-------|------------|---------|
| Name | EmpID | DeptName | Name | EmpID | DeptName | Mgr |
| Harry | 3415 | Finance | Harry | 3415 | Finance | NULL |
| Sally | 2241 | Sales | Sally | 2241 | Sales | Harriet |
| George | 3401 | Finance | George | 3401 | Finance | NULL |
| Harriet | 2202 | Sales | Harriet | 2202 | Sales | Harriet |
| | | | NULL | NULL | Production | Charles |

| Dept | |
|------------|---------|
| DeptName | Mgr |
| Sales | Harriet |
| Production | Charles |

Relational algebra Notation

| | | | | | |
|--------------|--------------|--------------|-------------------|------------------|--------------------------|
| Projection | PROJECT | π | Cartesian product | \times | \times |
| Selection | SELECT | σ | Join | JOIN | \bowtie |
| Renaming | RENAME | ρ | Left outer join | LEFT OUTER JOIN | \bowtie_{left} |
| Union | UNION | \cup | Right outer join | RIGHT OUTER JOIN | \bowtie_{right} |
| Intersection | INTERSECTION | \cap | Full outer join | FULL OUTER JOIN | \bowtie_{full} |
| Assignment | \leftarrow | \leftarrow | Semijoin | SEMIJOIN | \bowtie_{semi} |

Relational Calculus

- Comes in two flavors: Tuple relational calculus (TRC) and Domain relational calculus (DRC).
- Calculus has *variables*, *constants*, *comparison ops*, *logical connectives* and *quantifiers*.
 - TRC:** Variables range over (i.e., get bound to) *tuples*.
 - DRC:** Variables range over *domain elements* (= field values).
 - Both TRC and DRC are simple subsets of first-order logic.
- Expressions in the calculus are called formulas. An answer tuple is essentially an

DRC Formulas

- Atomic formula:**
 - $\langle x_1, x_2, \dots, x_n \rangle \in R_{name}$, or $X \text{ op } Y$, or $X \text{ op constant}$ $\langle, >, =, \leq, \geq, \neq$
 - op** is one of
- Formula:**
 - an atomic formula, or
 - $\neg p, p \wedge q, p \vee q$
 - $\exists X (p(X))$, where p and q are formulas,
 - or $\forall X (p(X))$, where X is a *domain variable* or
 - $\forall X \exists Y (p(X, Y))$, where X is a *domain variable*.
- The use of **quantifiers** and is said to bind X.

Free and Bound Variables

- The use of **quantifiers** $\exists X$ and $\forall X$ in a formula is said to bind X.
 - A variable that is **not bound** is free.
- Let us revisit the definition of a **query**:

$$\{ \langle x_1, x_2, \dots, x_n \rangle \mid p(\langle x_1, x_2, \dots, x_n \rangle) \}$$
 - There is an important restriction: the variables x_1, \dots, x_n that appear to the left of \mid must be the *only* free variables in the formula $p(\dots)$.

Find sailors rated > 7 who've reserved a red boat

$$\{ \langle I, N, T, A \rangle \mid \langle I, N, T, A \rangle \in \text{Sailors} \wedge T > 7 \wedge \exists Ir, Br, D (\langle Ir, Br, D \rangle \in \text{Reserves} \wedge Ir = I \wedge \exists B, BN, C (\langle B, BN, C \rangle \in \text{Boats} \wedge B = Br \wedge C = 'red')) \}$$

- Observe how the parentheses control the scope of each quantifier's binding.
- This may look cumbersome, but with a good user interface, it could be intuitive. (MS Access, QBE)

Find sailors who've reserved all boats

$$\{ \langle I, N, T, A \rangle \mid \langle I, N, T, A \rangle \in \text{Sailors} \wedge \\ \forall \langle B, BN, C \rangle \left(\neg \left(\langle B, BN, C \rangle \in \text{Boats} \right) \vee \right. \\ \left. \left(\exists \langle Ir, Br, D \rangle \left(\langle Ir, Br, D \rangle \in \text{Reserves} \wedge I = Ir \wedge Br = B \right) \right) \right\}$$

- Find all sailors I such that for each B , tuple either it is not a tuple in Boats or there is a tuple in Reserves showing that sailor I has reserved it.

Find sailors who've reserved all boats (again!)

$$\{ \langle I, N, T, A \rangle \mid \langle I, N, T, A \rangle \in \text{Sailors} \wedge \\ \forall \langle B, BN, C \rangle \in \text{Boats} \\ \left(\exists \langle Ir, Br, D \rangle \in \text{Reserves} (I = Ir \wedge Br = B) \right) \}$$

- Simpler notation, same query. (Much clearer!)
- To find sailors who've reserved all red boats:

$$\dots \{ C \neq \text{'red'} \vee \exists \langle Ir, Br, D \rangle \in \text{Reserves} (I = Ir \wedge Br = B) \}$$

Any other way to specify it? Equivalence in logic

Unsafe Queries, Expressive Power

- It is possible to write syntactically correct calculus queries that have an infinite number of answers! Such queries are called *unsafe*.
 - e.g., $\{ S \mid \neg (S \in \text{Sailors}) \}$
- It is known that every query that can be expressed in relational algebra can be expressed as a safe query in DRC / TRC; the converse is also true.
- Relational Completeness:** Query language (e.g., SQL) can express every query that is expressible in relational algebra/calculus.

Tuple Relational Calculus

- Interested in finding tuples for which a predicate is true. Based on use of tuple variables.
- Tuple variable is a variable that 'ranges over' a named relation: i.e., variable whose only permitted values are tuples of the relation.
- Specify range of a tuple variable S as the Staff relation as:
Staff(S)
- To find set of all tuples S such that $P(S)$ is true:
 $\{ S \mid P(S) \}$

Tuple relational calculus

- Similar to DRC except that variables range over **tuples** rather than field values
- For example, the query “Find all sailors with rating above 7” is represented in TRC as follows:

$\{S \mid S \in \text{Sailors} \wedge S.\text{rating} > 7\}$

49

Example

Find names and ages of sailors with a rating above 7

$\{P \mid \exists S \in \text{Sailors}. S.\text{rating} > 7 \wedge$
P.sname=S.sname \wedge
P.age=S.age}

Recall P ranges
over tuple values

50

Example

Find the names of sailors who have reserved at least two boats

$\{ P \mid \exists S \in \text{Sailors}.$
 $\exists R_1 \in \text{Reserves}.$
 $\exists R_2 \in \text{Reserves}.$
 $S.\text{sid} = R_1.\text{sid} \wedge R_1.\text{sid} = R_2.\text{sid}$
 \wedge
 $R_1.\text{bid} \neq R_2.\text{bid} \wedge$
 $P.\text{sname} = S.\text{sname} \}$

51

Encoding relational calculus

- Can we code up the relational calculus in the relational algebra?
- At the moment, **NO!**
- Given our syntax we can define ‘problematic’ queries such as
 $\{S \mid \neg (S \in \text{Sailors})\}$
- This (presumably) means the set of all tuples that are not sailors, which is an infinite set... 😞

52

Summary

- The relational model has rigorously defined query languages that are simple and powerful.
- Relational algebra is more operational; useful as internal representation for query evaluation plans.
- Several ways of expressing a given query; a query optimizer should choose the most efficient version.